

Lecture 4 Notes: Arrays and Strings

1 Arrays

So far we have used variables to store values in memory for later reuse. We now explore a means to store *multiple* values together as one unit, the *array*.

An array is a fixed number of *elements* of the same type stored sequentially in memory. Therefore, an integer array holds some number of integers, a character array holds some number of characters, and so on. The size of the array is referred to as its *dimension*. To declare an array in C++, we write the following:

```
type arrayName[dimension];
```

To declare an integer array named `arr` of four elements, we write `int arr[4];`

The elements of an array can be accessed by using an *index* into the array. Arrays in C++ are zero-indexed, so the first element has an index of 0. So, to access the third element in `arr`, we write `arr[2];` The value returned can then be used just like any other integer.

Like normal variables, the elements of an array must be initialized before they can be used; otherwise we will almost certainly get unexpected results in our program. There are several ways to initialize the array. One way is to declare the array and then initialize some or all of the elements:

```
int arr[4];

arr[0] = 6;
arr[1] = 0;
arr[2] = 9;
arr[3] = 6;
```

Another way is to initialize some or all of the values at the time of declaration:

```
int arr[4] = { 6, 0, 9, 6 };
```

Sometimes it is more convenient to leave out the size of the array and let the compiler determine the array's size for us, based on how many elements we give it:

```
int arr[] = { 6, 0, 9, 6, 2, 0, 1, 1 };
```

Here, the compiler will create an integer array of dimension 8.

The array can also be initialized with values that are not known beforehand:

```
1  #include <iostream>
2  using namespace std;
3
4  int main() {
5
6      int arr[4];
7      cout << "Please enter 4 integers:" << endl;
8
9      for(int i = 0; i < 4; i++)
10         cin >> arr[i];
11
```

```

12     cout << "Values in array are now:";
13
14     for(int i = 0; i < 4; i++)
15         cout << " " << arr[i];
16
17     cout << endl;
18
19     return 0;
20 }

```

Note that when accessing an array the index given must be a positive integer from 0 to $n-1$, where n is the dimension of the array. The index itself may be directly provided, derived from a variable, or computed from an expression:

```

arr[5];
arr[i];
arr[i+3];

```

Arrays can also be passed as arguments to functions. When declaring the function, simply specify the array as a parameter, without a dimension. The array can then be used as normal within the function. For example:

```

0  #include <iostream>
1  using namespace std;
2
3  int sum(const int array[], const int length) {
4      long sum = 0;
5      for(int i = 0; i < length; sum += array[i++]);
6      return sum;
7  }
8
9  int main() {
10     int arr[] = {1, 2, 3, 4, 5, 6, 7};
11     cout << "Sum: " << sum(arr, 7) << endl;
12     return 0;
13 }

```

The function `sum` takes a constant integer array and a constant integer length as its arguments and adds up `length` elements in the array. It then returns the sum, and the program prints out Sum: 28.

It is important to note that arrays are *passed by reference* and so any changes made to the array within the function will be observed in the calling scope.

C++ also supports the creation of multidimensional arrays, through the addition of more than one set of brackets. Thus, a two-dimensional array may be created by the following:

```

type arrayName[dimension1][dimension2];

```

The array will have *dimension1* x *dimension2* elements of the same type and can be thought of as an array of arrays. The first index indicates which of *dimension1* subarrays to access, and then the second index accesses one of *dimension2* elements within that subarray. Initialization and access thus work similarly to the one-dimensional case:

```

1  #include <iostream>
2  using namespace std;
3
4  int main() {
5      int twoDimArray[2][4];
6      twoDimArray[0][0] = 6;
7      twoDimArray[0][1] = 0;
8      twoDimArray[0][2] = 9;

```

```

9     twoDimArray[0][3] = 6;
10    twoDimArray[1][0] = 2;
11    twoDimArray[1][1] = 0;
12    twoDimArray[1][2] = 1;
13    twoDimArray[1][3] = 1;
14
15    for(int i = 0; i < 2; i++)
16        for(int j = 0; j < 4; j++)
17            cout << twoDimArray[i][j];
18
19    cout << endl;
20    return 0;
21 }

```

The array can also be initialized at declaration in the following ways:

```

int twoDimArray[2][4] = { 6, 0, 9, 6, 2, 0, 1, 1 };

int twoDimArray[2][4] = { { 6, 0, 9, 6 } , { 2, 0, 1, 1 } };

```

Note that dimensions must *always* be provided when initializing multidimensional arrays, as it is otherwise impossible for the compiler to determine what the intended element partitioning is. For the same reason, when multidimensional arrays are specified as arguments to functions, all dimensions but the first *must* be provided (the first dimension is optional), as in the following:

```

int aFunction(int arr[][4]) { ... }

```

Multidimensional arrays are merely an abstraction for programmers, as all of the elements in the array are sequential in memory. Declaring `int arr[2][4];` is the same thing as declaring `int arr[8];`.

2 Strings

String literals such as "Hello, world!" are actually represented by C++ as a sequence of characters in memory. In other words, a string is simply a character array and can be manipulated as such.

Consider the following program:

```

1  #include <iostream>
2  using namespace std;
3
4  int main() {
5      char helloworld[] = { 'H', 'e', 'l', 'l', 'o', ',', ' ', 'w',
6                          'o', 'r', 'l', 'd', '!', '\0' };
7
8      cout << helloworld << endl;
9
10     return 0;
11 }

```

This program prints `Hello, world!` Note that the character array `helloworld` ends with a special character known as the *null character*. This character is used to indicate the end of the string.

Character arrays can also be initialized using string literals. In this case, no null character is needed, as the compiler will automatically insert one:

```

char helloworld[] = "Hello, world!";

```

The individual characters in a string can be manipulated either directly by the programmer or by using special functions provided by the C/C++ libraries. These can be included in a program through the use of the `#include` directive. Of particular note are the following:

- `cctype` (`cctype.h`): character handling
- `cstdio` (`stdio.h`): input/output operations
- `cstdlib` (`stdlib.h`): general utilities
- `cstring` (`string.h`): string manipulation

Here is an example to illustrate the `cctype` library:

```
1  #include <iostream>
2  #include <cctype>
3  using namespace std;
4
5  int main() {
6      char messyString[] = "t6H0I9s6.iS.999a9.STRING";
7
8      char current = messyString[0];
9      for(int i = 0; current != '\0'; current = messyString[++i]) {
10         if(isalpha(current))
11             cout << (char)(isupper(current) ? tolower(current) : current);
12         else if(ispunct(current))
13             cout << ' ';
14     }
15
16     cout << endl;
17     return 0;
18 }
```

This example uses the `isalpha`, `isupper`, `ispunct`, and `tolower` functions from the `cctype` library. The `is-` functions check whether a given character is an alphabetic character, an uppercase letter, or a punctuation character, respectively. These functions return a Boolean value of either `true` or `false`. The `tolower` function converts a given character to lowercase.

The for loop beginning at line 9 takes each successive character from `messyString` until it reaches the null character. On each iteration, if the current character is alphabetic and uppercase, it is converted to lowercase and then displayed. If it is already lowercase it is simply displayed. If the character is a punctuation mark, a space is displayed. All other characters are ignored. The resulting output is `this is a string`. For now, ignore the `(char)` on line 11; we will cover that in a later lecture.

Here is an example to illustrate the `cstring` library:

```
1  #include <iostream>
2  #include <cstring>
3  using namespace std;
4
5  int main() {
6      char fragment1[] = "I'm a s";
7      char fragment2[] = "tring!";
8      char fragment3[20];
9      char finalString[20] = "";
10
11     strcpy(fragment3, fragment1);
12     strcat(finalString, fragment3);
13     strcat(finalString, fragment2);
14
15     cout << finalString;
16     return 0;
17 }
```

This example creates and initializes two strings, `fragment1` and `fragment2`. `fragment3` is declared but not initialized. `finalString` is partially initialized (with just the null character).

`fragment1` is copied into `fragment3` using `strcpy`, in effect initializing `fragment3` to `I'm a s.` `strcat` is then used to concatenate `fragment3` onto `finalString` (the function overwrites the existing null character), thereby giving `finalString` the same contents as `fragment3`. Then `strcat` is used again to concatenate `fragment2` onto `finalString`. `finalString` is displayed, giving `I'm a string!`.

You are encouraged to read the documentation on these and any other libraries of interest to learn what they can do and how to use a particular function properly. (One source is <http://www.cplusplus.com/reference/>.)